



VAPTOR AI



VAPTOR AI

Autonomous, AI-powered penetration testing and vulnerability assessment
— running continuously on your network, not once a year in a consultant's
report.

PLUG-IN APPLIANCE · CONTINUOUS SCANNING · VERIFIED EXPLOITS · AI RISK ANALYSIS

Security testing hasn't kept pace with risk

Most organizations still validate their security posture the way they did a decade ago: an annual or twice-yearly penetration test, scoped weeks in advance, delivered as a static PDF months after the testing window closed. In between engagements, networks change, new services go live, and freshly disclosed CVEs sit unverified — leaving long, unmonitored windows of exposure that attackers don't wait for.

Skilled penetration testers are expensive and difficult to schedule on demand, and compliance frameworks are moving from "show us a report" to "show us continuous assurance." Stretching a lean security team to keep up with both is no longer realistic.

THE VAPTOR AI SOLUTION

A self-contained platform that thinks and tests like a pentester

VAPTOR AI is a turnkey appliance and AI-driven testing platform that runs the same playbook a professional penetration tester would — discovery, vulnerability scanning, exploit verification, and AI-prioritized risk analysis — on whatever schedule your organization needs. Results land in a live dashboard and an evidence-ready report your team and auditors can act on immediately, not months later.

ALWAYS-ON ASSURANCE

Scheduled scans run daily, weekly, or monthly, so new exposures are caught between audits — not at the next one.

AI-PRIORITIZED FINDINGS

Every finding is enriched and explained in plain language, so your team fixes what matters first.

AUDIT-READY EVIDENCE

Executive summaries, technical detail, and compliance mappings are generated automatically, every run.

New vulnerabilities published per year

Worldwide CVE publications, 2023–2025

Year	Publications (k)
2023	28.8k
2024	40.0k
2025	48.2k

~133 new CVEs are published every day in 2025 — a 263% increase in annual submissions since 2020. Source: CVE.org / Jerry Gamblin 2025 CVE Data Review.

How attackers are getting in

Share of breaches by initial access vector, 2025

Initial Access Vector	Share of Breaches
Vulnerability exploitation	20%
Phishing	15%

Vulnerability exploitation as a breach entry point grew 34% year-over-year and has now overtaken phishing. Exploitation of edge devices and VPNs alone grew nearly 8x (3% → 22%), and only ~54% of those flaws were fully remediated, in a median of 32 days. Source: Verizon 2025 Data Breach Investigations Report.

54.8 days

**AVG. TIME TO REMEDIATE A
CRITICAL VULNERABILITY**

Edgescan 2025 Vulnerability Statistics Report

45.4%
**OF ENTERPRISE VULNERABILITIES
STILL UNPATCHED AFTER 12
MONTHS**

Edgescan 2025 Vulnerability Statistics Report


~40%
**OF ORGANIZATIONS TEST ONLY
ONCE OR TWICE A YEAR**
2024–25 industry pentest-frequency surveys

Speed of detection is the single biggest cost lever

Breach cost isn't driven mainly by sophistication of the attacker — it's driven by how long the attacker goes unnoticed. That's the lever a continuously-running, AI-prioritized platform is built to pull, every day instead of twice a year.

Average breach cost by containment speed

Global average, 2025



The 241-day average breach lifecycle in 2025 was the fastest in nine years — yet most breaches still cross the 200-day mark where costs jump roughly 30%. Source: IBM Cost of a Data Breach Report 2025.

\$4.44M

GLOBAL AVERAGE COST OF A DATA BREACH (2025)

IBM Cost of a Data Breach Report 2025

\$1.9M saved

AVERAGE SAVINGS FOR ORGANIZATIONS USING AI EXTENSIVELY IN DETECTION & RESPONSE — PLUS AN 80-DAY-SHORTER BREACH LIFECYCLE

IBM Cost of a Data Breach Report 2025

\$10.5T

PROJECTED GLOBAL COST OF CYBERCRIME IN 2025

Cybersecurity Ventures

The math, in plain terms

A traditional pentest engagement runs roughly \$10,000–\$35,000 and covers a single point in time, once or twice a year — leaving the other 350+ days uninstrumented. VAPTOR AI runs the same discovery-to-exploit-verification pipeline on a schedule you set, with AI-driven prioritization on every finding, so the gap between "tested" and "today" stops being measured in months.

Figures above are independently published industry benchmarks (cited per chart) describing the broader threat and cost landscape. They illustrate the stakes continuous testing is built to address and are not a guarantee of savings for any specific deployment.

From discovery to decision — automatically

Every VAPTOR AI engagement runs through the same disciplined six-phase pipeline a senior penetration tester would follow — without needing to schedule one.

1

Discovery

Network and port scanning maps every live host and service in scope.

2

Scouting

Deep reconnaissance enumerates technology stacks, web applications, and exposed services.

3

Pentest

Web, SQL, and exploit-verification testing (OWASP ZAP, SQLMap, Metasploit) confirms what's actually exploitable.

4

Assessment

AI enrichment cross-references findings against live CVE and threat intelligence feeds and scores risk.

5


Report

Executive summary and technical detail are generated automatically — re-runnable, never hand-edited.


6

Complete

Findings land in the dashboard with remediation tracking, ready for your team and auditors.



Live dashboard — active scans, total vulnerabilities, verified exploits, asset count, and an aggregate risk score, updated in real time as each phase completes.



At-a-glance fleet health — scans running, vulnerabilities found, and exploits verified across every monitored target.

One platform, the full assessment toolkit

Industry-standard engines (Nmap, Masscan, OWASP ZAP, Nuclei) identify exposed services and web application weaknesses.

Metasploit-driven exploit verification confirms which vulnerabilities are genuinely exploitable — not just theoretically present.

Automated SQLMap-based testing of web inputs and parameters for injection flaws.



Live traffic capture flags credential exposure and unsafe protocols on the wire — without ever storing a real password.

Suricata-based monitoring evaluates perimeter detection coverage during a scan.

Every finding is enriched with CVE context and summarized in plain-English, board-ready language.

Read-only agents pull Windows, Linux, and macOS security logs and surface findings automatically.


Daily, weekly, or monthly scans with quick, full, and stealth intensity profiles to match your environment.



Built to the standard it's testing for

A tool that probes your network for weaknesses has to be held to a higher security standard than the systems it tests. VAPTOR AI is engineered accordingly.

- ✓ **No plaintext credentials, ever.** Passwords are hashed with PBKDF2-SHA256 at 260,000 iterations with a unique random salt per user.
- ✓ **Role-based access control.** Four roles — Admin, Operator, Analyst, Client — govern exactly who can launch scans, run exploits, view reports, or manage settings.
- ✓ **JWT-secured web access** with short-lived, signed sessions and server-side permission checks on every request.
- ✓ **TLS-verified, audited API calls.** Every outbound integration is encrypted and logged; sensitive fields are automatically redacted.
- ✓ **Credentials are never harvested.** Network-observed credentials are masked, hashed, and never stored or displayed in raw form.
- ✓ **Hardware-bound licensing.** A device's license key is permanently bound to its hardware, so a lost or copied key can't be reused elsewhere.
- ✓ **Parameterized everything.** All database access uses parameterized queries — no string-built SQL, anywhere in the platform.



Role-based permission matrix — administrators control, at a glance, exactly what each role can see and do — and can update it live without a restart.

Plug it in. It starts testing.


VAPTOR AI ships as a self-contained hardware appliance — the VAPTOR Watchdog — that connects directly to your network. There's no lengthy professional-services install, no agents to roll out across endpoints, and no dependency on a third party's cloud to run a scan.



VAPTOR Watchdog — a compact, plug-and-play appliance that brings the full testing platform onto your network in minutes.

Each device carries its own identity, bound permanently to its hardware at first registration — so device management stays simple even as your footprint grows, and a decommissioned device can be revoked instantly and can't be reused elsewhere.

Multi-organization support is built in from the ground up: assets, scans, findings, and report history are isolated per organization, so a single VAPTOR AI deployment can safely serve multiple business units or multiple clients at once.



Centralized fleet management — for MSSPs and consultancies running VAPTOR AI across multiple client sites, the Control Panel tracks every device, license, and update from one place — usage, license tier, and status, all in one view.

A tier for every stage of your security program

TIER	WHAT'S INCLUDED	SCAN CAPACITY
Starter	Device + Adhoc Support	20 Assets
Professional	Device + Support (SLA)	≤100 Scan Assets
Enterprise	Device + Management + Support (SLA)	>100 Scan Assets

Every tier includes the VAPTOR Watchdog appliance, the full scan pipeline, AI-driven risk analysis, and reporting. Contact us for tier pricing and a deployment plan suited to your environment.

See VAPTOR AI running on your own network

Request a demo and we'll walk through a live scan, a sample executive report, and how VAPTOR AI fits into your existing security and compliance workflow.

sales@vaptgpt.com cyberwatchdogs.co.za vaptgpt.com